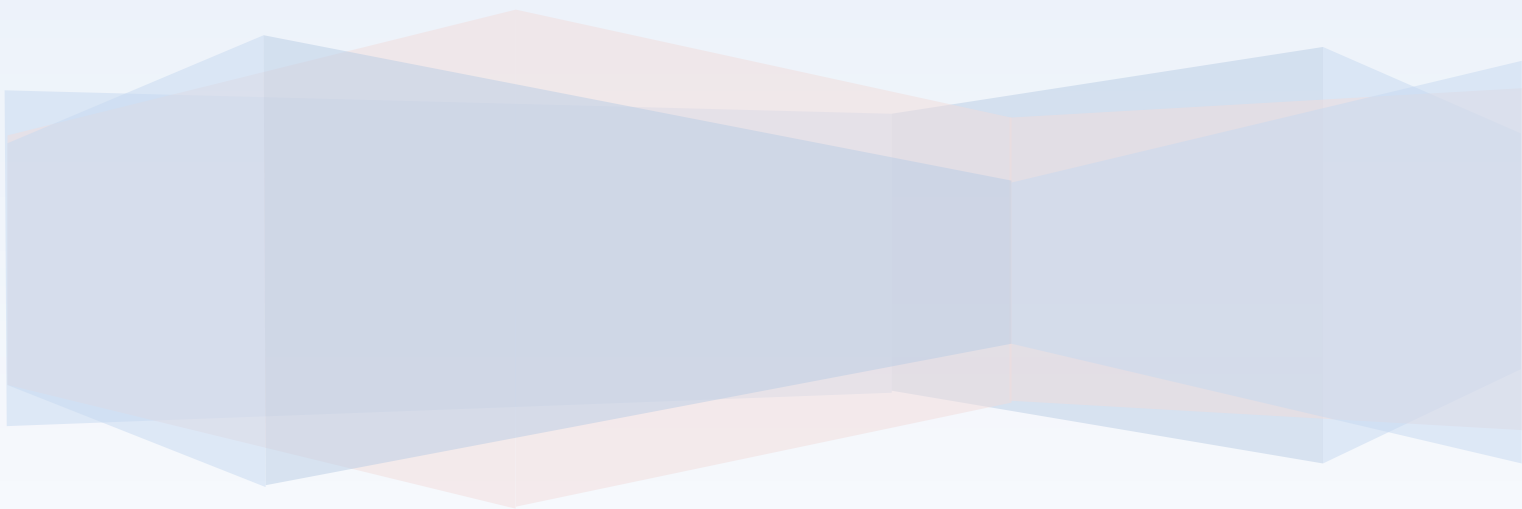


COS30018 Intelligent Systems Assignment B

Semester 2, 2023

Assignemnt B Report

Robin Findlay-Marks (103603871) and the Project leader: Mukesh Malani



Self-Assessment Details

The following checklists provide an overview of my self-assessment for this unit.

	Pass (P)	Credit (C)	Distinction (D)	High Distinction (Low HD) (High HD)	
Self-Assessment (please tick)	✓				

Self-assessment Statement

	Included? (tick)
Learning Summary Report	✓
Complete Pass ("core") task work, approved in Canvas	✓

Minimum Pass Checklist

	Included? (tick)
Additional non-core task work (or equivalent) in a private repository and accessible to staff account.	
Spike Extension Report (for spike extensions) in Canvas	
Custom Project plan (for D and/or low HD), and/or High HD Research Plan document in Canvas (optional)	

Credit Checklist, in addition to Pass Checklist

Introduction

This report summarises what I have done for the COS30018 Intelligent Systems Assignment B. For this semester I was given a basic stock prediction file as a starting point (v0.1). Over the following weeks I was given several tasks to incrementally improve the v0.1 code until it is done.

Instructions to run the program

To run the program, first install the requirements listed in the requirements.txt file which are as follows:

```
sklearn
tensorflow
matplotlib
numpy
pandas
pandas-datareader
yahoo_fin
yfinance
scikit-learn
mplfinance
statsmodels
cython
prophet
```

These can be installed individually with the pip install command, or by running the pip install command on the requirements.txt file.

The next step to run the program is to open the Intelligent-Systems-Assignemnt-B folder in the command prompt.

Finally run the command to run the stock_prediction.py file which is:
 py stock_prediction.py

To change the settings of the program, open the parameters.py, which has all the settings and parameters used in the program.

One important setting to note is the 'MODE' parameter. This parameter decides which of several modes to run the program in. They are; 1 – to use the basic RNN models with the data split into train and test by a specified date, 2 - to use the basic RNN models with the data split into train and test by a ratio, 3 – display candlestick chart of the past NDAYS of downloaded data, 4 – display boxplot chart of the past NDAYS of downloaded data, 5 – run the prophet prediction model and any other number, which uses the basic RNN models with the data split into train and test by a random date.

Some additional functionality instructions (assumes parameters are set to default for each instruction)

To test the Multivariate function, set the MODE parameter to 2 and MULTIVARIATE parameter to true

To test the Ensemble with ARIMA function, set the MODE parameter to 2 and ENSEMBLE parameter to true

To test the Ensemble with SARIMA function, set the MODE parameter to 2, ENSEMBLE parameter to true and SARIMA parameter to true

To test the Random Forest function, set the MODE parameter to 2, FOREST parameter to true

To test the hyperparameters, scroll down on the ADVANCED SETTINGS and then Hyperparameters section in parameters.py.

To change the RNN model used, change the hyperparameters's RNN_HYPERPARAM to 1 for LSTM, 2 for SimpleRNN, 3 for GRU

Overall system architecture

The overall program uses just 2 python files. The stock_prediction.py file runs all the functions, while the parameters.py file has all the settings and parameters for the program.

The parameters.py file has various settings such as the ticker symbol, start and end date of the stock data range and how many days into the future to predict are present. Each setting should have a comment to explain what it does and I have already explained what is needed to run the different modes of the program in section of this document detailing how to run the program. As such I will not explain how parameters.py will work.

For the stock_prediction.py file, it starts with the Main() function at the bottom of the file. This function takes the input from the MODE parameter and puts it into a switch to determine what functions and functionality to use.

Main function used:

```

558 def Main(): #Main function for deciding what functions to use
559     #Make filename for the saved data file
560     ticker_data_filename = os.path.join("data", f"{COMPANY}_{TRAIN_START}_{TEST_END}.csv")
561
562     # Switch for checking which mode to run the program in
563     match MODE:
564         case 1: #Predict with date split
565             getDataSplitDate(ticker_data_filename, SPLIT_DATE)
566             runTest()
567
568         case 2: #Predict with ratio split
569             getDataRatio(ticker_data_filename, RATIO)
570             runTest()
571
572         case 3: #Candlestick Chart
573             candlestickChart(ticker_data_filename)
574
575         case 4: #Boxplot Chart
576             boxplotChart(ticker_data_filename)
577
578         case 5: #Prophet Prediction
579             getDataRatio(ticker_data_filename, RATIO)
580             runTestProphet()
581
582         case _: #Predict with random date split
583             #Convert dates to datetime
584             dateStart = datetime.strptime(TRAIN_START, '%Y-%m-%d')
585             dateEnd = datetime.strptime(TEST_END, '%Y-%m-%d')
586             #Get random date inbetween the start and end
587             random_date = dateStart + (dateEnd - dateStart) * random.random()
588             #Convert back to string
589             random_date = random_date.strftime('%Y-%m-%d')
590             #Use random date as split
591             getDataSplitDate(ticker_data_filename, random_date)
592             runTest()
593
594     Main()

```

Implemented data processing techniques

The program implements several data processing techniques. This starts with the checkFiles function, which can be seen below.

```

37 # Function for checking if the data is already downloaded (needs internet connection to download)
38 # If the data is NOT in a file, it downloads the data and makes a csv file and returns the data
39 # If the data IS in a file, it reads the data from the file and returns the data
40 def checkFiles(filename):
41     if (os.path.exists(filename)):
42         #Read csv file and return the data inside
43         data = pd.read_csv(filename)
44         # NaN values from pandas are values that are not a number. For example in stocks if the stock data for a
45         # specific day was not recorded, I believe it would still have a record for that day, only the values
46         # would be NaN or 'Not a Number'
47
48         #what dropna() does is simply remove the missing values from the dataset
49         data.dropna(inplace=True)
50         return data
51     else:
52         #Download data from online
53         data = yf.download(COMPANY, start=TRAIN_START, end=TEST_END, progress=False)
54
55         # Save data to csv file
56         data.to_csv(filename)
57         # For some reason it needs to read it from the file otherwise it won't work
58         data = pd.read_csv(filename)
59
60         if (STOREFILE == False):
61             os.remove(filename) #Remove stored file
62         # remove NaN values from the dataset
63         data.dropna(inplace=True)
64         return data

```

This function checks if the wanted data has already been downloaded. If it has been, it loads the data from the file, otherwise it downloads the data from the internet and then stores it as a file.

The next step is to split the data into training and test datasets, which is run in one of three ways. The first, is to split the data by a specified date, the code can be seen below,

```

66 # This function gets the datafile name as well as the split date
67 # it then runs the file checker to get the dataset, then splits the dataset at the split date
68 # and sets the trainData and testData
69 def getDataSplitDate(filename, splitDate):
70     df = checkFiles(filename)
71
72     # Make it know that the date column is indeed a date
73     df['Date'] = pd.to_datetime(df['Date'])
74     df = df.set_index(df['Date'])
75     df = df.sort_index()
76
77     #Convert input to datetime, add 1 day, then convert back to string
78     date = datetime.strptime(splitDate, '%Y-%m-%d')
79     testStartDate = date + timedelta(days=1)
80     testStartDate = testStartDate.strftime('%Y-%m-%d')
81
82     # Set the fullData variable
83     global fullData
84     fullData = df
85
86     # create train/test partition
87     global trainData
88     trainData = df[TRAIN_START:splitDate]
89     trainData = trainData.drop(trainData.columns[[0]], axis=1)
90     global testData
91     testData = df[testStartDate:TEST_END]
92     testData = testData.drop(testData.columns[[0]], axis=1)
93     print('Train Dataset:', trainData.shape)
94     print('Test Dataset:', testData.shape)

```

the second splits the data by a ratio, the code can be seen below,

```

96 # This function gets the datafile name as well as the 'ratio' number
97 # it then runs the file checker to get the dataset, then splits the dataset at the split date
98 # and sets the trainData and testData
99 def getDataRatio(filename, ratio):
100     df = checkFiles(filename)
101
102     # Make it know that the date column is indeed a date
103     df['Date'] = pd.to_datetime(df['Date'])
104     df = df.set_index(df['Date'])
105     df = df.sort_index()
106
107     # Convert strings to dates
108     date1 = datetime.strptime(TRAIN_START, '%Y-%m-%d')
109     date2 = datetime.strptime(TEST_END, '%Y-%m-%d')
110
111     # do math to get the date we want
112     trainEndDate = date2 + (date1 - date2) / ratio
113
114     #Convert input to datetime, add 1 day
115     print("Middle : " + trainEndDate.strftime('%Y-%m-%d'))
116     testStartDate = trainEndDate + timedelta(days=1)
117     # Convert back to string
118     testStartDate = testStartDate.strftime('%Y-%m-%d')
119     trainEndDate = trainEndDate.strftime('%Y-%m-%d')
120
121     # Code for testing dates
122     print('trainEndDate Dataset:', trainEndDate)
123     print('testStartDate:', testStartDate)
124
125     # Set the fullData variable
126     global fullData
127     fullData = df
128
129     # create train/test partition
130     global trainData
131     trainData = df[TRAIN_START:trainEndDate]
132     trainData = trainData.drop(trainData.columns[[0]], axis=1)
133     global testData
134     testData = df[testStartDate:TEST_END]
135     testData = testData.drop(testData.columns[[0]], axis=1)

```

and the uses a random date to split it, by generating a random date between the start and end date, then inputting that into the split by date function, which can be seen below.

```

582         case _: #Predict with random date split
583             #Convert dates to datetime
584             dateStart = datetime.strptime(TRAIN_START, '%Y-%m-%d')
585             dateEnd = datetime.strptime(TEST_END, '%Y-%m-%d')
586             #Get random date inbetween the start and end
587             random_date = dateStart + (dateEnd - dateStart) * random.random()
588             #convert back to string
589             random_date = random_date.strftime('%Y-%m-%d')
590             #Use random date as split
591             getDataSplitDate(ticker_data_filename, random_date)
592             runTest()
593

```

Another data processing method that is present is to display the downloaded data as a candlestick chart (code below)

```

523 def candlestickChart(filename):
524     # Get data
525     df = checkFiles(filename)
526
527     # Make it know that the date column is indeed a date
528     df['Date'] = pd.to_datetime(df['Date'])
529     # Set the index of the dataframe to be the date column
530     df = df.set_index(df['Date'])
531     df = df.sort_index()
532
533     # Get the last n days
534     actual_prices_small = df[-NDAYS:]
535
536     # Plot the candlestick chart
537     mpl.plot(actual_prices_small.set_index("Date"), type="candle", style="charles", title='Candlestick Chart')
538     # Uses mplfinance to make the candlestick chart
539     # uses the date column as index
540     # uses the 'charles' style to make the decreasing days red and increasing days green

```

and a boxplot chart (code below).

```

542 def boxplotChart(filename):
543     # Get data
544     df = checkFiles(filename)
545     # Get the last n days
546     actual_prices_small = df[-NDAYS:]
547
548     # Plot the boxplot chart
549     fig = actual_prices_small[['Open', 'Close', 'Low', 'High']].plot(kind='box', title='Boxplot Chart', grid=True)
550     # Uses Matplotlib to make the boxplot chart
551     # uses grid=True to make a grid so the chart is easier to read
552
553     # Add axis labels and display the plot
554     plt.xlabel("Stock Data type")
555     plt.ylabel("Price")
556     plt.show()
557

```

I believe that these are all the main data processing techniques used in the program.

Implemented machine learning techniques

The program implements several machine learning techniques. The first of these is a function to function with several inputs and return a deep learning model. The input parameters are the number of layers, the size of each layer and the RNN model subtype. The code for this function can be seen below.

```

243 def createModelRNN(layer_num, layer_size, layer_name, dropout):
244     #Declare some variables so the model knows what's what
245     PRICE_VALUE = "close"
246
247     scaler = MinMaxScaler(feature_range=(0, 1))
248     scaled_data = scaler.fit_transform(trainData[PRICE_VALUE].values.reshape(-1, 1))
249
250     # To store the training data
251     x_train = []
252     y_train = []
253
254     scaled_data = scaled_data[:,0] # Turn the 2D array back to a 1D array
255     # Prepare the data
256     for x in range(LOOKBACK_DAYS, len(scaled_data)):
257         x_train.append(scaled_data[x-LOOKBACK_DAYS:x])
258         y_train.append(scaled_data[x])
259
260     # Convert them into an array
261     x_train, y_train = np.array(x_train), np.array(y_train)
262     # Now, x_train is a 2D array(p,q) where p = len(scaled_data) - LOOKBACK_DAYS
263     # and q = LOOKBACK_DAYS; while y_train is a 1D array(p)
264     x_train = np.reshape(x_train, (x_train.shape[0], x_train.shape[1], 1))
265     # We now reshape x_train into a 3D array(p, q, 1); Note that x_train
266     # is an array of p inputs with each input being a 2D array

```

```

267 model = Sequential() # Basic neural network
268 # Add layers to network using for each loop, which takes the layer_num to determine how many layers are added
269 for i in range(layer_num):
270     if i == 0:
271         # first layer
272         model.add(layer_name(layer_size, return_sequences=True, input_shape=(x_train.shape[1], 1)))
273     elif i == layer_num - 1:
274         # last layer
275         model.add(layer_name(layer_size, return_sequences=False))
276     else:
277         # hidden layers
278         model.add(layer_name(layer_size, return_sequences=True))
279         # add dropout after each layer
280         model.add(dropout(dropout))
281
282 # Prediction of the next closing value of the stock price
283 model.add(Dense(units=1))
284 # Compile the model
285 model.compile(optimizer='adam', loss='mean_squared_error')
286 # Now we are going to train this model with our training data
287 # (x_train, y_train)
288 model.fit(x_train, y_train, epochs=25, batch_size=32)
289
290 # Return completed model to be tested
291 return model

```

The input variables that are used for this function are modified from the parameters.py, with various hyperparameter presets in the file. They can be viewed below.

```

63 # Default Parameters
64
65 #RNN param
66 LAYER_NUM = 2 # number of layers used
67 LAYER_SIZE = 50 # size of each layer
68 LAYER_NAME = LSTM # which type of RNN model is used
69 DROPOUT = 0.2 # dropout

```

```

84 RNN_HYPERPARAM = 1
85 ARIMA_HYPERPARAM = 1
86 SARIMA_HYPERPARAM = 1
87
88 match RNN_HYPERPARAM:
89     case 1: #LSTM
90         LAYER_NUM = 2
91         LAYER_SIZE = 50
92         LAYER_NAME = LSTM
93         DROPOUT = 0.2
94     case 2: # RNN
95         LAYER_NUM = 2
96         LAYER_SIZE = 50
97         LAYER_NAME = SimpleRNN
98         DROPOUT = 0.2
99     case 3: #GRU
100         LAYER_NUM = 2
101         LAYER_SIZE = 50
102         LAYER_NAME = GRU
103         DROPOUT = 0.2
104     case 4: #P1 settings
105         LAYER_NUM = 2
106         LAYER_SIZE = 256
107         LAYER_NAME = LSTM
108         DROPOUT = 0.4
109     case 5: #Custom
110         LAYER_NUM = 2
111         LAYER_SIZE = 50
112         LAYER_NAME = GRU
113         DROPOUT = 0.1
114     case _: #Default settings
115         LAYER_NUM = 2
116         LAYER_SIZE = 50
117         LAYER_NAME = SimpleRNN
118         DROPOUT = 0.2

```

The next machine learning technique used is the multistep prediction. This means that the program must be able to predict an amount of days into the future, determined by an input. For this I made the RNN model predict a day in the future, then add that to the model's input data, which is then used to make a prediction another day into the future. This process repeats, until the input (PREDICTION_DAYS) days into the future is predicted. The code for this part of the function is below.

```

448 futurePrice = []
449
450
451 i = 0
452
453 # for each day in the future to predict
454 while i < PREDICTION_DAYS:
455     i += 1
456
457     # make it so it does (or redoes) the declaring of real data based on the last PREDICTION_DAYS amount of days
458     real_data = [model_inputs[len(model_inputs) - PREDICTION_DAYS:, 0]]
459     real_data = np.array(real_data)
460     real_data = np.reshape(real_data, (real_data.shape[0], real_data.shape[1], 1))
461
462     # make prediction of next day
463     prediction = model.predict(real_data)
464     # unscale and flatten prediction data
465     scaledPrediction = scaler.inverse_transform(prediction)
466     # add predicted data to future price
467     futurePrice.append(scaledPrediction.flatten()[0])
468     # set prediction to be the flattened but still scaled data
469     prediction = prediction.flatten()[0]
470
471     print(futurePrice)
472
473     # set model inputs to be dataframe, add predicted data to it, then make it a numpy array again
474     model_inputs = pd.DataFrame(model_inputs)
475     model_inputs.loc['0'] = prediction
476     model_inputs = model_inputs.to_numpy()
477

```

The next machine learning technique used is the multivariate prediction. This means that the program must use the data other than just the closing data as inputs to predict the closing price.

The code for this can be viewed below

```

167 def multivariate_prediction(layer_num, layer_size, layer_name):
168     # make a copy of the train and test dataframes
169     train_df = train_data.sort_values(by=['date']).copy()
170     test_df = test_data.sort_values(by=['date']).copy()
171     # Add dummy column and set dummy values for scaling in the future
172     train_df_ext = train_df.copy()
173     train_df_ext['dummy'] = train_df_ext['close']
174     test_df_ext = test_df.copy()
175     test_df_ext['dummy'] = test_df_ext['close']
176     # Get the number of rows in the data
177     n_rows = train_df.shape[0]
178     # Convert the data to numpy values
179     np_train_unscaled = np.array(train_df)
180     np_test_unscaled = np.array(test_df)
181     np_data = np.reshape(np_train_unscaled, (n_rows, -1))
182     # Transform the data by scaling each feature to a range between 0 and 1
183     scaler = MinMaxScaler()
184     np_train_scaled = scaler.fit_transform(np_train_unscaled)
185     np_test_scaled = scaler.fit_transform(np_test_unscaled)
186     # Creating a separate scaler that works on a single column for scaling predictions
187     scaler_pred = MinMaxScaler()
188     df_close = pd.DataFrame(train_df_ext['close'])
189     df_close2 = pd.DataFrame(test_df_ext['close'])
190     np_close_scaled = scaler_pred.fit_transform(df_close)
191     np_close_scaled2 = scaler_pred.fit_transform(df_close2)
192     # Set Prediction Index
193     index_close = train_df.columns.get_loc("close")
194     # Create the training and test data
195     train_data = np_train_scaled
196     test_data = np_test_scaled
197     # Here, we create N samples, LOOKBACK_DAYS time steps per sample, and 6 features
228     # Prediction of the next closing value of the stock price
229     model.add(Dense(1))
230     # Compile the model
231     model.compile(optimizer='adam', loss='mean_squared_error')
232     # Training the model
233     early_stop = EarlyStopping(monitor='loss', patience=5, verbose=1)
234     history = model.fit(x_train, y_train, batch_size=32, epochs=10, validation_data=(x_test, y_test))
235     # Get the predicted values
236     y_pred_scaled = model.predict(x_test)
237     # Unscale the predicted values
238     y_pred = scaler_pred.inverse_transform(y_pred_scaled)
239     y_test_unscaled = scaler_pred.inverse_transform(y_test.reshape(-1, 1))
240
241     return y_pred

```

```

198 # Here, we create N samples, LOOKBACK_DAYS time steps per sample, and 6 features
199
200 def partition_dataset(LOOKBACK_DAYS, data):
201     x, y = [], []
202     data_len = data.shape[0]
203     for i in range(LOOKBACK_DAYS, data_len):
204         x.append(data[i-LOOKBACK_DAYS:i,:]) #contains LOOKBACK_DAYS values 0-LOOKBACK_DAYS * columns
205         y.append(data[i, index_close]) #contains the prediction values for validation, for single-step prediction
206     # Convert x and y to numpy arrays
207     x = np.array(x)
208     y = np.array(y)
209     return x, y
210
211 # Generate training data and test data
212 x_train, y_train = partition_dataset(LOOKBACK_DAYS, train_data)
213 x_test, y_test = partition_dataset(LOOKBACK_DAYS, test_data)
214
215 # Configure the neural network model
216 model = Sequential()
217
218 # Add layers to network using for each loop, which takes the layer_num to determine how many layers are added
219 for i in range(layer_num):
220     if i == 0:
221         # first layer
222         model.add(layer_name(layer_size, return_sequences=True, input_shape=(x_train.shape[1], x_train.shape[2])))
223     elif i == layer_num - 1:
224         # last layer
225         model.add(layer_name(layer_size, return_sequences=False))
226     else:
227         # hidden layers
228         model.add(layer_name(layer_size, return_sequences=True))
229
230 # Prediction of the next closing value of the stock price

```

The final machine learning technique used that is not part of the extension is an ensemble model of the RNN model and a ARIMA or SARIMA model. An option to use ARIMA or SARIMA is present. The ensemble approach I took, was to get the values for each model and get the average of them to get the ensemble value.

The S/ARIMA code can be viewed below

```

137 # Function for running ARIMA or SARIMA predictions
138 def ARIMA_prediction():
139     # assign train and test data to variables
140     train = train_data['close'].values
141     test1 = test_data[1:]
142     test = test1['close'].values
143     history = [x for x in train]
144     predictions = list()
145     # parameters for SARIMA
146     my_seasonal_order = (SAUTOREG, SDIFFERENCE, SMOVAVG, SEASON)
147
148     # walk-forward validation
149     for t in range(len(test)):
150         # re-create the ARIMA model after each new observation
151         if SARIMA:
152             model = ARIMA(history, order=(AUTOREG, DIFFERENCE, MOVAVG), seasonal_order=my_seasonal_order)
153         else:
154             model = ARIMA(history, order=(AUTOREG, DIFFERENCE, MOVAVG))
155         model_fit = model.fit()
156         # make prediction
157         output = model_fit.forecast()
158         forecast = output[0]
159         predictions.append(forecast)
160         expected = test[t]
161         # keep track of past observations
162         history.append(expected)
163         print('predicted=%f, expected=%f' % (forecast, expected))
164     return predictions

```

The ensembling of the data code can be viewed below

```

495 # combine the RNN prediction value and the S/ARIMA prediction value and average them to make an ensemble value
496 ensemble_preds = None
497 if ENSEMBLE:
498     predicted_prices_list = predicted_prices.tolist()
499     i = 0
500     # for each predicted value from each mode, add the two together and to make the ensemble value
501     for value in arima_pred:
502         # first parameter is the array that the second parameter is being added to
503         ensemble_preds = np.append(ensemble_preds, (arima_pred[i] + predicted_prices_list[i])/2)
504         i += 1

```

Details of extension work

For the extension work I decided to add two new machine learning techniques. The first of these is using facebook's prophet to make predictions. The code for this is fairly simple concise, being only one function, however due to it's unique data input method, it has it's own separate MODE from the rest of the code.

```

426 def runTestProphet():
427     # get pre-split data
428     # Only need to extract Close because the date is the index, and as such is automatically transferred over too
429     test = fullData['Close']
430     # reset index because it turns it back into a normal column which is referenced later
431     test = test.reset_index()
432     # turn date and close column into ds and y (necessary for prophet to work)
433     test.columns = ['ds', 'y']
434     # make sure ds column is a datetime
435     test['ds'] = pd.to_datetime(test['ds'])
436     # remove offset days from the training data
437     train = test.drop(test.index[-PROPHET_TRAIN_OFFSET:])
438
439     # make the prophet model
440     model = Prophet()
441     # train the model from the train data
442     model.fit(train)
443     # setup dataframe from only the date date
444     futuredays = list()
445     futuredays = test['ds']
446     futuredays = pd.DataFrame(futuredays)
447
448     # use the model to make a forecast from the futuredays datetime range
449     forecast = model.predict(futuredays)
450     # set the actual and predicted values
451     actualData = test['y'][-len(forecast):].values
452     prophetPrediction = forecast['yhat'].values
453     # plot actual vs Predicted Data
454     plt.plot(actualData, label='Actual')
455     plt.plot(prophetPrediction, label='Predicted')
456     plt.legend()
457     plt.show()

```

The code for this can be viewed below

The other machine learning technique I added was random forest. This was more complicated having a main function, and three sub functions. The code for this can be viewed below.

```

410 def runTestForest():
411     # get test + train data and make local variables
412     global testData
413     global trainData
414     train = trainData["Close"].values
415     test = testData["Close"].values
416
417     # turn train and test data into supervised learning
418     train = data_to_supervised(train)
419     test = data_to_supervised(test)
420
421     # make predictions
422     forestPrediction = forest_validation(test, train)
423
424     return forestPrediction
363 # turn time series data into a supervised learning data
364 def data_to_supervised(data):
365     df = pd.DataFrame(data)
366     columns = list()
367
368     # sliding window technique used to make the new samples for the supervised learning data
369     # input sequence (t-n, ... t-1)
370     for i in range(1, 0, -1):
371         columns.append(df.shift(i))
372     # forecast sequence (t, t+1, ... t+n)
373     for i in range(0, 1):
374         columns.append(df.shift(-i))
375     # put it all together
376     newdf = concat(columns, axis=1)
377     # drop NaN values
378     newdf.dropna(inplace=True)
379     return newdf.values
381 # fit an random forest model and make a one step prediction
382 def forest_prediction(train, testx):
383     # turn list into an array
384     train = asarray(train)
385     # split into input and output columns
386     trainX, trainy = train[:, :-1], train[:, -1]
387     # make model
388     model = RandomForestRegressor(n_estimators=FOREST_ESTIMATORS)
389     model.fit(trainX, trainy)
390     # make a single prediction
391     forestPrediction = model.predict([testx])
392     return forestPrediction[0]

```

```


# loop for running the prediction on a number of days
def forest_loop(test, train):
    predictions = list()
    # seed history with training dataset
    history = [x for x in train]
    # step over each time-step in the test set
    for i in range(len(test)):
        # split test row into input and output columns
        testx = test[i, :-1]
        # fit model on history and make a prediction
        forestPrediction = forest_prediction(history, testx)
        # store forecast in list of predictions
        predictions.append(forestPrediction)
    return predictions

```


Due to the small page limit for this document, please refer to the task 7 report for more info on how these techniques were implemented

Demonstration of working system

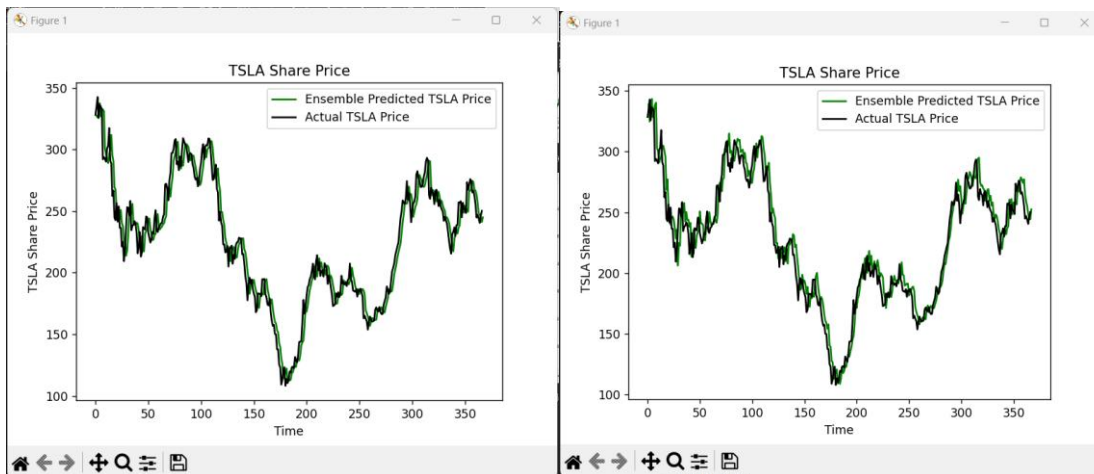
RNN model predictions with LSTM, then SimpleRNN, then GRU



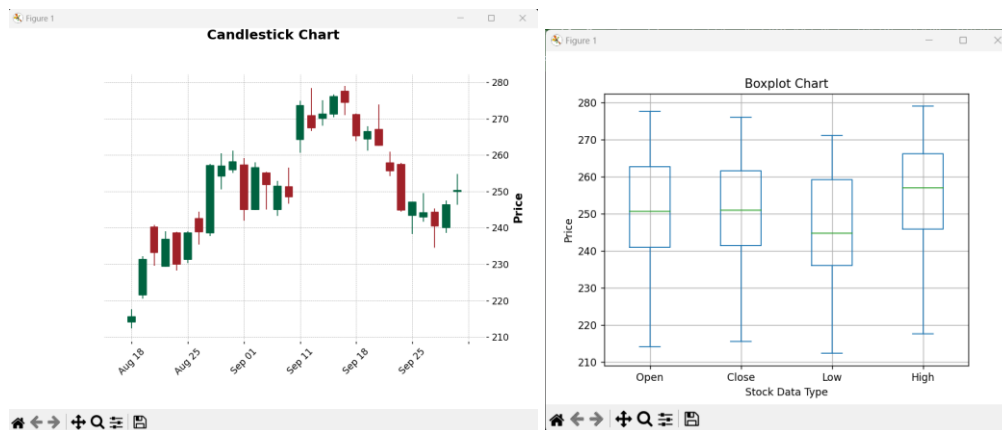
First image is 60 days in the future prediction (multistep), then second is Multivariate prediction



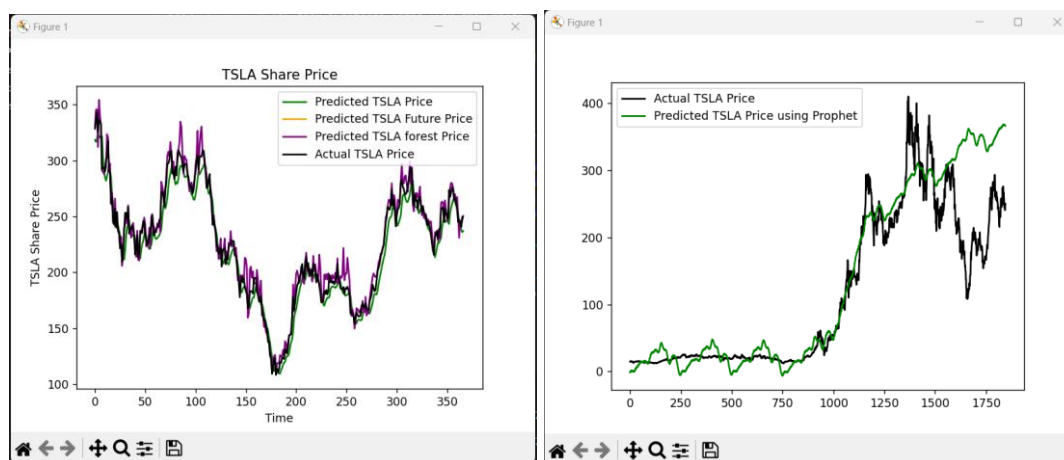
First image is Arima ensemble predictions, then Sarima ensemble predictions



First image is candlestick chart, then boxplot chart



First image is random forest predictions, then the prophet predictions



Reflection on work done

Overall I am happy with the work I did for this assignment. I felt like I learnt a lot about the basics for how to make and run several different machine learning models and what each model is useful for. If in the future I use machine learning models I will have already working code that I made that I can use as a launching pad to get the work done.

I would have liked to be able to get the prophet predictions working so they make more than vaguely similar predictions to the actual data. I am happy with how the random forest code runs, although it clearly needs some tweaking to get that working as well as some of the other models.

My only worry is that I completed this report in the way that was expected of me. I am unsure if I was meant to include this many screenshots of the work I did or if it was meant to be filled with explanations of how the code works. I concluded that it should be screenshots and a summary since reports 1-7 contained more detailed explanations of the work done. Please refer to these reports if you wish for further explanations.

Summary/Conclusion

For this assignment B, I spent a lot of hours working to get all the features of the program working and I believe I succeeded. In addition to this, I also added some extra machine learning functionality in the form of the simple extension.

URL to github:

<https://github.com/R0binicus/Intellegent-Systems-Assigment-B>